


LLM SOTA Pricing Analysis — May 2026

May 20, 2026

LLM SOTA Pricing Analysis - May 2026



Executive Summary

The LLM API market in mid-2026 is defined by one word: **asymmetry**. Chinese labs (Kimi, Z.AI/GLM, MiniMax) are systematically undercutting Western counterparts by 60–99% while matching or exceeding benchmark performance. American and European providers have responded with moderate price cuts from 2023–2024 peaks, but their premium tiers remain priced for enterprise contracts with unlimited budgets.

Key insight: The best-performing model per dollar is not from OpenAI, Anthropic, or Google. It is from a Chinese startup selling tokens at prices that look like typos.

The "Are You Sure This Isn't a Typo?" Tier

Model	Input	Output	Context	Source
DeepSeek V4 Flash	\$0.14	\$0.28	-	DeepSeek

DeepSeek V4 Pro	\$1.74	\$3.48	-	DeepSeek
Grok 4.1 Fast	\$0.20	\$0.50	2M	xAI
Gemini 2.5 Flash	\$0.175	\$0.75	-	Google
Gemini 3.1 Flash-Lite	\$0.25	\$1.50	-	Google
GLM-4-32B-0414	\$0.10	\$0.10	128K	Z.AI
GLM-4.7-FlashX	\$0.07	\$0.40	-	Z.AI
GLM-4.7-Flash	FREE	FREE	-	Z.AI
MiniMax M2.5	\$0.15	\$1.15	205K	MiniMax
MiniMax M2.1	\$0.29	\$0.95	205K	MiniMax
MiniMax M1	\$0.40	\$2.20	1M	MiniMax

The "Reasonable CFOs" Tier

Model	Input	Output	Context	Source
GPT-4o mini	\$0.15	\$0.60	-	OpenAI
GPT-5 nano	\$0.05	\$0.40	-	OpenAI
Gemini 3 Flash	\$0.50	\$3.00	-	Google
Claude Haiku 4.5	\$1.00	\$5.00	-	Anthropic
Mistral Medium 3	\$0.40	\$2.00	-	Mistral
Kimi K2.5	\$0.40	\$1.90	262K	Moonshot AI
GLM-4.7	\$0.60	\$2.20	200K	Z.AI
GLM-4.5-Air	\$0.20	\$1.10	-	Z.AI
MiniMax-01	\$0.20	\$1.10	-	MiniMax
MiniMax M2.7	\$0.279	\$1.20	205K	MiniMax

The "We're Serious About AI" Tier (Flagships)

Model	Input	Output	Context	Source
GPT-5.2	\$1.75	\$14.00	400K	OpenAI
Gemini 2.5 Pro	\$1.25	\$10.00	1M	Google
Gemini 3.1 Pro	\$2.00	\$12.00	1M	Google
Claude Sonnet 4.6	\$3.00	\$15.00	200K	Anthropic

GLM-5 / GLM-5.1	\$1.00–\$1.40	\$3.20–\$4.40	200K	Z.AI
Kimi K2.6	\$0.73	\$3.49	262K	Moonshot AI
MiniMax M2.5-Lightning	\$0.30	\$2.40	205K	MiniMax

The "Expensed to a VC Firm" Tier

Model	Input	Output	Context	Source
Claude Opus 4.6 / 4.7	\$5.00	\$25.00	200K	Anthropic
GPT-5.4	\$2.50	\$15.00	-	OpenAI
GPT-5.5	\$5.00	\$30.00	-	OpenAI
GPT-5.4 Pro	\$30.00	\$180.00	1.05M	OpenAI
GLM-4.5-X	\$2.20	\$8.90	-	Z.AI

The "Your Board Will Ask Questions" Tier

Model	Input	Output	Source
Claude Mythos Preview	\$25.00	\$125.00	Anthropic
GPT-5.2 Pro	\$21.00	\$168.00	OpenAI
o1-pro	\$150.00	\$600.00	OpenAI

Chinese Labs: The Disruption Detail

Kimi (Moonshot AI)

- **Kimi K2.6:** 1T parameter open-weight model. \$0.73/\$3.49. Beats GPT-5.4 on SWE-Bench Pro. 262K context.
- **Kimi K2.5:** \$0.40/\$1.90. Visual coding specialist. Previous generation still highly competitive.
- Philosophy: Growth over gross margin. \$20B valuation, noodle-stall pricing.

Z.AI / GLM (Zhupu AI)

- **GLM-5.1:** \$1.40/\$4.40. Latest flagship. Beats Claude Opus 4.6 on coding.
- **GLM-5:** \$1.00/\$3.20. Previous flagship.
- **GLM-4.7-FlashX:** \$0.07/\$0.40. Near-free high-speed tier.
- **GLM-4-32B-0414:** \$0.10/\$0.10. Both directions. Disposable pricing.

- **GLM-4.7-Flash / GLM-4.5-Flash:** Actually free. Limited but real.
- Philosophy: Permanent sale. Cache inputs for literal pennies (\$0.01–\$0.26/M).

MiniMax

- **MiniMax M2.7:** \$0.279/\$1.20. Self-improvement + agentic orchestration.
- **MiniMax M2.5:** \$0.15/\$1.15. 80.2% SWE-Bench Verified. Sonnet 4.6 territory for gumball money.
- **MiniMax M1:** \$0.40/\$2.20 with **1M context**. Compare to Claude \$5/\$25 or GPT-5.4 Pro \$30/\$180.
- Also offers full creative suite: video, speech, music at prices that make Runway and ElevenLabs look like luxury brands.
- Philosophy: "If it costs more than a coffee, we failed."

Value Scorecard

Lab	Best Value Pick	Score-Per-Dollar Hero
DeepSeek	V4 Flash at \$0.14/\$0.28	260.7 ratio - undisputed champion
Z.AI / GLM	GLM-4-32B at \$0.10/\$0.10	Free Flash models exist
MiniMax	M2.5 at \$0.15/\$1.15	Best agentic coding under \$1
Kimi	K2.5 at \$0.40/\$1.90	Long-context SOTA under \$2
Google	Gemini 3.1 Pro at \$2/\$12	Best Western flagship value
OpenAI	GPT-5.2 at \$1.75/\$14	Most improved from GPT-4 era
Anthropic	Claude Sonnet 4.6 at \$3/\$15	The "sane person" pick

Consumer Subscriptions (Chat Access)

Service	Monthly Price
ChatGPT Plus	\$20
Claude Pro	\$20 (~\$17 annual)
X Premium+ (Grok)	\$22
Gemini Enterprise	\$30/user
Claude Max	\$100–\$200
SuperGrok Heavy	\$300

Open-Weight Self-Hosting Note

Models like Llama 4, Qwen 3.5, Nemotron, Kimi K2.6, and MiniMax M2.5 are available as open weights.

Download cost: **\$0**. Inference cost on cloud GPUs: roughly **\$0.50-\$5.00/hour** depending on model size and quantization. "Free" is doing heavy lifting, but at scale it can undercut even Chinese API pricing.

Conclusion

If you are building on US or EU APIs at list price in May 2026, you are either burning VC money or you have not checked the exchange rate on Chinese tokens. The performance gap has closed. The price gap has not. Choose accordingly.

Compiled May 20, 2026. Sources: BenchLM.ai, IntuitionLabs.ai, CloudIDR.com, TheNextWeb.com, Z.AI official docs, OpenRouter, TokenMix.ai, MiniMax.io.